

Creating value with automation

3.1 Democratizing Machine Learning

3.2 Model Factories

3.3 Continuous Learning

Democratizing Machine Learning

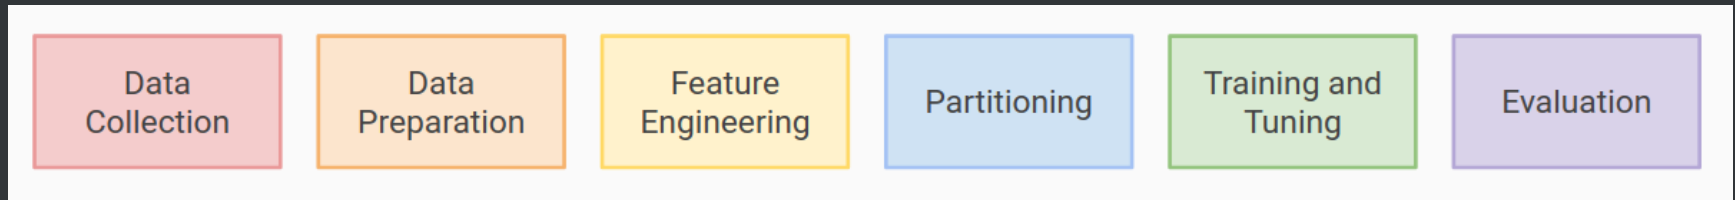
If AI is indeed *the new electricity* - how can we ensure that everyone gets access to it?

Barriers for democratization?

- Workforce shortage
 - Data Scientists cost top dollars and often lack the domain expertise required.
 - Empower Domain experts (e.g. Business Analysts, SWE) instead?
- Research -> Production Gap
 - Building models is easy - getting them into production and business processes is hard.
 - ~80% of ML don't get into production.
 - Unicorn: ML Engineer (strong Data Scientist and Software Engineer w/ DevOps skills; impossible to find).
 - [MLOps](<https://en.wikipedia.org/wiki/MLOps>) to the rescue \o/
- Access to data
 - Data sources
 - Data catalogues

True End-to-End Learning

- Assist/ automate every step in the ML Process



- Open Problems
 - Problem formulation
 - Translate Business Problems into an objective that can be optimized.
 - Select metric and model selection scheme
- Data acquisition & cleansing
 - Automated Feature Engineering
- Operation & maintenance (more on that later).

How can AML help in Democratization?

Empower Business Analysts, Data Engineers, SWE, and Citizen Scientists to use ML to solve problems by providing capabilities to

- build accurate ML models reliably with little to no human involvement (data acquisition, metric/partitioning selection, pipeline opt),
- with guardrails to avoid pitfalls and catastrophic failure (leakage),
- and guidance when the model should be used and when not!

Where does the AML community need to step up:

- Everywhere but pipeline optimization and hyper-parameter tuning

Model Factories

As organizations become more data driven, **reliable analytics and data science** will become an essential part of staying competitive and keeping costs under control.

Many Data Science teams, however, still **develop models in an ad-hoc fashion** on their workstations and hand over trained models to Data Engineers or SWEs for productionalization.

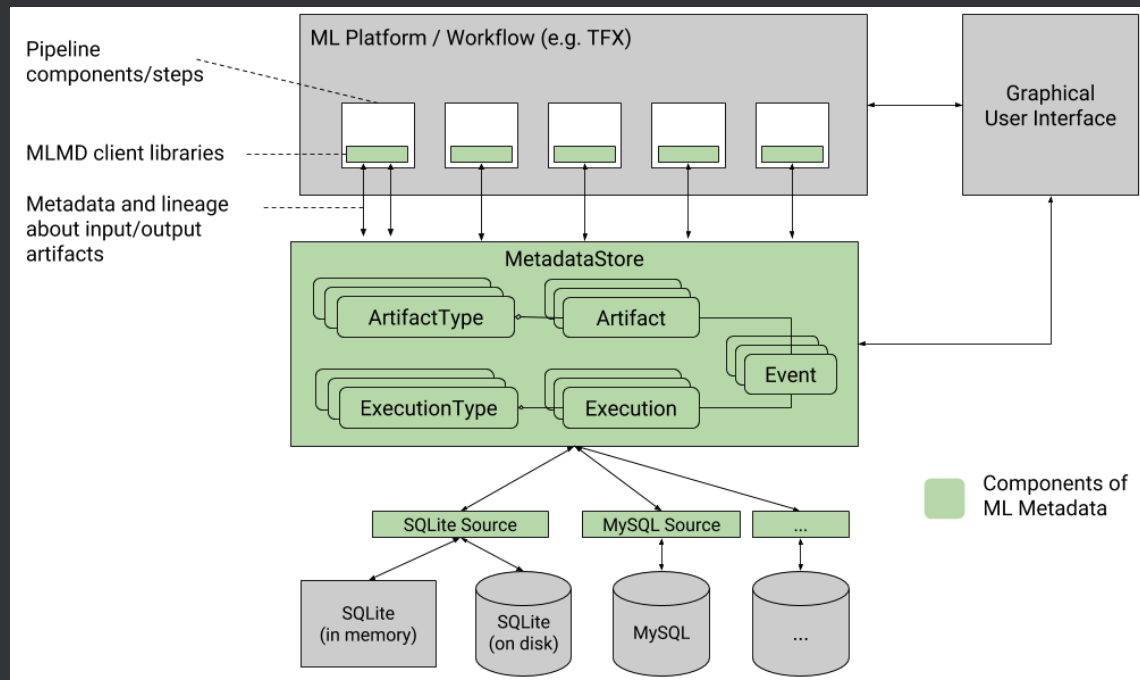
Requirements

- Version control of models & archiving
- Automated build & testing
- Input data checks
- Reproducibility & Lineage
- Governance & regulatory compliance
- Monitoring
- Scheduling

Model Metadata Stores

MLMD

- Provides Lineage
 - Reproducibility
 - Checkpointing (Pause/Resume)
- Versioning



How can AML help in Model Factories?

Consistent Quality

Having a solid model selection and assessment framework minimizes the surface for (human) error.

Governance

Platform ensures that lineage is tracked and metadata recorded (who built what model when and how).

Case-Study: Kubeflow

Open source ecosystem for Machine Learning (automation) for Kubernetes.

Not a *Model Factory* but contains relevant building blocks.

Ecosystem components:

- **Pipelines**: Workflow orchestration
- **Argo CD**: Continuous Delivery, GitOps
- **Fairing**: package models trained in a Jupyter notebook
- **KFServing** / SeldonCore: Model Deployment / Serving
- **Katib**: Hyper-parameter tuning



Continuous Learning

The world constantly changes... this begs the question:

- Are my model assessment results still valid?
- Am I doing worse... or can I do better?

ML Models make assumptions about the data generating process, we need to automatically recognize changes in data.

Data Drift

Sources of changes

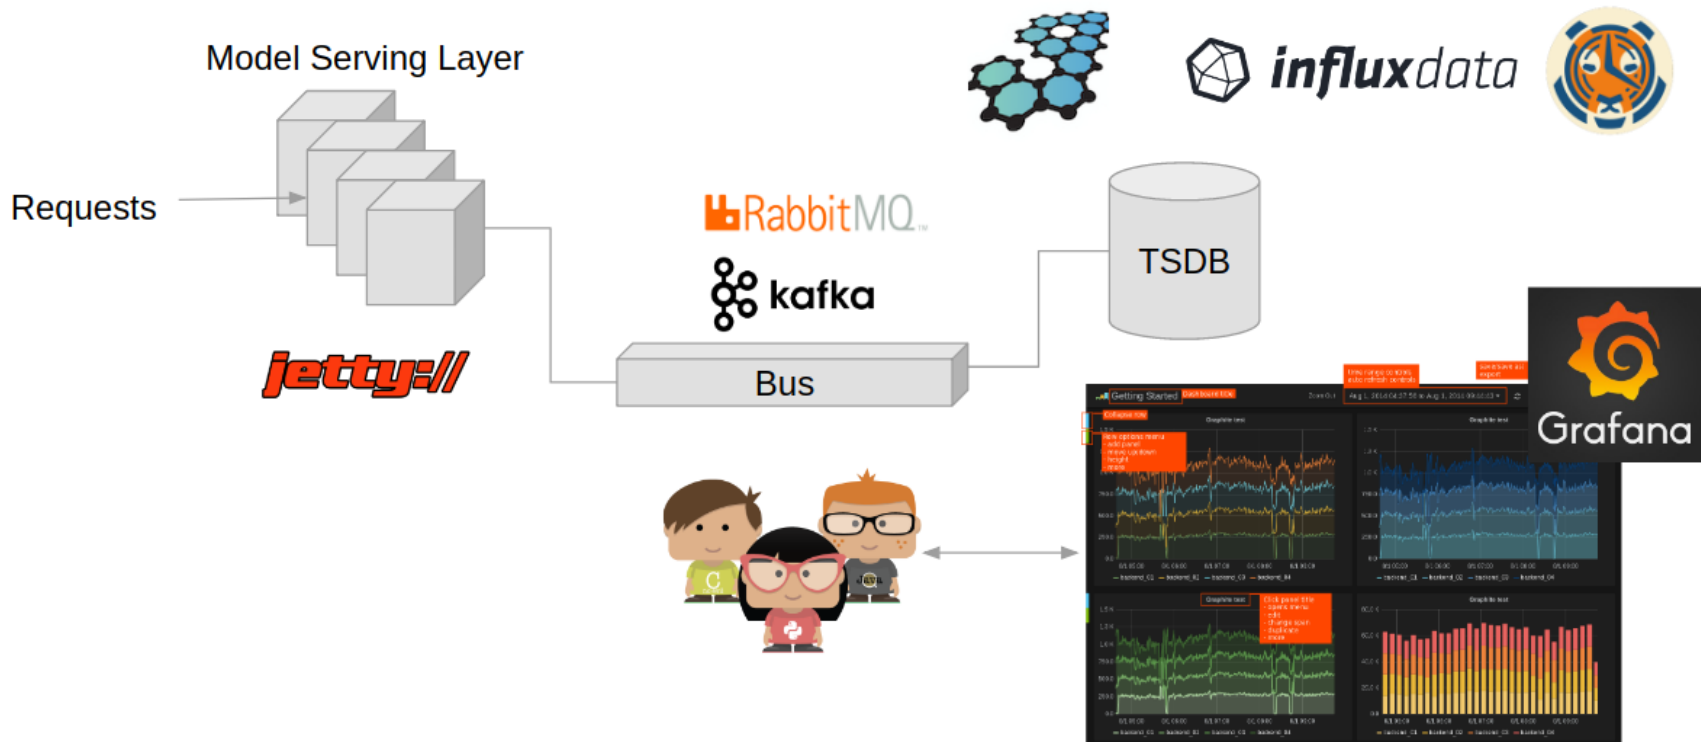
- Changes in DB design, broken sensors, new semantics, new user-interface alters interaction pattern, ...
- In real world, DBs and data sources are living, breathing, evolving entities

Automatically detect drift

- When labels are available: run model assessment again and compare to old values.
- When labels are not available
 - Look for changes in the distribution of the predictions (using histogram distance metrics like Population-Stability-Index)
 - Look for simple univariate statistic
 - Fraction of missing values
 - Fraction of new categorical levels
 - Fraction of values outside a certain range (e.g. $2 * \text{std}$)

Automatically detecting drift is hard; regularly scheduled retrains are more common.

Model Monitoring Architecture



What to do when Data Drift is detected?

- Warning flags / alert
 - the model might still work (as measured by other KPIs)
- Default to a more robust model
 - This is where AML can help: quickly build a model without a drifting feature.
- Re-train model
 - This is usually a lengthy process..
- Adapt existing model
 - This is for a different time...



Case-study: TFX



TensorFlow Extended (TFX) is an end-to-end platform for deploying production ML pipelines

Key building blocks (not exhaustive):

- TF Data Validation: Understand, validate and monitor data.
- TF Transform: Feature preprocessing as a TF graph; limits what you can do but provides safety and interoperability.
- **ML Metadata**: Integral part of TFX, ensures compatibility of different artifacts (models)

TFX Training Orchestration

Airflow DAGs Data Profiling Browse Admin Docs About 2019-07-28 15:55:03 UTC

On DAG: taxi schedule: None

Graph View Tree View Task Duration Task Tries Landing Times Gantt Details Code Trigger DAG Refresh Delete

success Base date: 2019-07-11 12:31:00 Number of runs: 25 Run: manual__2019-07-11T12:30:59.148821+00:00 Layout: Left->Right Go Search for...

Component success running failed skipped up_for_reschedule up_for_retry queued no_status

```
graph LR; CsvExampleGen --> StatisticsGen; CsvExampleGen --> SchemaGen; CsvExampleGen --> ExampleValidator; StatisticsGen --> SchemaGen; SchemaGen --> Transform; SchemaGen --> Trainer; Transform --> Trainer; ExampleValidator --> Trainer; Trainer --> Evaluator; Trainer --> ModelValidator; Evaluator --> Pusher; ModelValidator --> Pusher;
```

The diagram illustrates the TFX Training pipeline. It starts with 'CsvExampleGen' which feeds into 'StatisticsGen', 'SchemaGen', and 'ExampleValidator'. 'StatisticsGen' also feeds into 'SchemaGen'. 'SchemaGen' feeds into 'Transform' and 'Trainer'. 'Transform' also feeds into 'Trainer'. 'ExampleValidator' also feeds into 'Trainer'. 'Trainer' feeds into 'Evaluator' and 'ModelValidator'. 'Evaluator' also feeds into 'Pusher'. 'ModelValidator' also feeds into 'Pusher'.